

Terraform Apply:

```
MINGW64:/c/Users/YuriiiHordiichuk/Desktop/m06_sparkbasics_python_aws-m... - X
module.eks.module.eks_managed_node_group["one"].aws_eks_node_group.this[0]: still creating... [01m50s elapsed]
module.eks.module.eks_managed_node_group["two"].aws_eks_node_group.this[0]: Creation complete after 1m51s [id=eks-g85pk5LY:node-group-2-202511171412542243000001b]
module.eks.module.eks_managed_node_group["one"].aws_eks_node_group.this[0]: still creating... [02m00s elapsed]
module.eks.module.eks_managed_node_group["one"].aws_eks_node_group.this[0]: Creation complete after 2m2s [id=eks-g85pk5LY:node-group-1-202511171412542237000009]

Apply complete! Resources: 68 added, 0 changed, 0 destroyed.

Outputs:
cluster_endpoint = "https://6FBA74FD2D0B83737D4CCB24983F0832.gr7.us-west-1.eks.amazonaws.com"
cluster_name = "eks-g85pk5LY"
cluster_security_group_id = "sg-0ca184c09971432cb"
region = "us-west-1"
```

Authenticate Docker to ECR:

```
BV1VEQkFFdU1CRUVESn1SQmE4cU1PQTNybhZtd3dJQkvJQTdLb01PaXJueFVCYmo3REpwVmhxMzRMM1ovWEJaQm15MHpUREM0d11EL3hjaGhSYmVpUmVPRHNGejQ1YmV0QWJ6cWQ5bGxocDgrb29QaWg0PSIsInZ1cnNpb24i0iIyIiwidHlwZSI6IkRBVEFs0VZIiwiZXhwaXJhdGlvbiI6MTc2MzQzMzAzMn0= 838240021663.dkr.ecr.us-west-1.amazonaws.com
WARNING! Using --password via the CLI is insecure. Use --password-stdin.
Login Succeeded
```

```
AzureAD\YuriiiHordiichuk@EPPL-WROW0218_MTNGW64 ~/Desktop/m06_sparkbasics_python_aw
```

Local tests result:

report.html

Report generated on 16-Nov-2025 at 15:24:48 by [pytest-html](#) v4.1.1

Environment

Python	3.11.0
Platform	Windows-10-10.0.26200-SP0
Packages	<ul style="list-style-type: none">• pytest: 9.0.1• pluggy: 1.6.0
Plugins	<ul style="list-style-type: none">• html: 4.1.1• metadata: 3.1.1
JAVA_HOME	C:\User\Program Files\Java\jdk-17.0.2

Summary

13 tests took 00:00:07.

(Un)check the boxes to filter the results.

0 Failed, 13 Passed, 0 Skipped, 0 Expected failures, 0 Unexpected passes, 0 Errors, 0 Reruns

[Show all details](#) / [Hide all details](#)

Result	Test	Duration	Links
Passed	src/test/etl_tests.py::TestGeocodingAPI::test_geocode_nyc_address	828 ms	
Passed	src/test/etl_tests.py::TestGeocodingAPI::test_geocode_london_address	590 ms	
Passed	src/test/etl_tests.py::TestGeocodingAPI::test_geocode_tokyo_address	374 ms	
Passed	src/test/etl_tests.py::TestLatLngWithAPI::test_get_lat_lng_with_real_api	437 ms	
Passed	src/test/etl_tests.py::TestLatLngWithAPI::test_get_lat_lng_london_with_real_api	393 ms	
Passed	src/test/etl_tests.py::TestGeohash::test_generates_geohash_from_valid_coordinates	1 ms	
Passed	src/test/etl_tests.py::TestGeohash::test_returns_none_for_null_latitude	0 ms	
Passed	src/test/etl_tests.py::TestGeohash::test_returns_none_for_both_null	0 ms	
Passed	src/test/etl_tests.py::TestGeohash::test_handles_invalid_input_gracefully	0 ms	
Passed	src/test/etl_tests.py::TestGeohash::test_different_locations_produce_different_hashes	1 ms	
Passed	src/test/etl_tests.py::TestEncryption::test_encrypts_specified_pii_fields	00:00:04	
Passed	src/test/etl_tests.py::TestEncryption::test_raises_error_for_invalid_fields	552 ms	

Docker build, tag, push:

```
MINGW64:/c/Users/YuriiiHordiichuk/Desktop/m06_sparkbasics_python_aws-m... ━ ━ ━ ×
=> CACHED [stage-0  6/15] RUN pip3 install --no-cache-dir --upgrade pip  0.0s
=> CACHED [stage-0  7/15] COPY ./entrypoint.sh /opt/                0.0s
=> [stage-0  8/15] COPY ./dist/sparkbasics-*.egg /opt/                0.0s
=> [stage-0  9/15] COPY --from=extra-source ./src /opt/src             0.0s
=> [stage-0 10/15] RUN chmod +x /opt/*.sh                            0.3s
=> [stage-0 11/15] RUN ln -s /usr/bin/tini /sbin/tini                 0.3s
=> [stage-0 12/15] RUN groupadd -g 999 spark && useradd -m -u 185 -g 0.3s
=> [stage-0 13/15] WORKDIR /opt/spark/work-dir                      0.1s
=> [stage-0 14/15] RUN mkdir -p /opt/spark/work-dir && chown -R spar 0.2s
=> [stage-0 15/15] RUN chown spark:spark /opt/sparkbasics-*.egg && chmod 0.3s
=> exporting to image                                                 0.1s
=> => exporting layers                                                0.1s
=> => writing image sha256:9046262460926755ba33b7ba962528dc5bd59085e6ec8 0.0s
=> => naming to docker.io/library/m06-spark-image                     0.0s
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ docker tag m06-spark-image:latest 838240021663.dkr.ecr.us-west-1.amazonaws.com
/spark-python-06:latest
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ docker push 838240021663.dkr.ecr.us-west-1.amazonaws.com/spark-python-06:lates
t
The push refers to repository [838240021663.dkr.ecr.us-west-1.amazonaws.com/spar
k-python-06]
3e78a17574c9: Pushed
5491a2c83852: Pushed
0181a1c1dcfc: Pushed
5b75bd9b9b9d: Pushed
eeb5fc5d4231: Pushed
f596e771c530: Pushed
86170511e037: Pushed
d88664949ebd: Pushed
7b4762837fdc: Pushed
3f0fc864f267: Pushed
6bbfe40fe5d5: Pushed
eddb715c2785: Pushed
39b2ff005b05: Pushed
6686d8a0d20c: Pushed
73974f74b436: Pushed
latest: digest: sha256:ae223e62d9fd517b13e7025683a390e24ec02162f75485c580b655dbf
92dfb06 size: 3454
(.venv)
```

Kubernetes Service Account Setup for Spark:

```
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ aws eks list-clusters --region us-west-1
{
    "clusters": [
        "eks-g85pk5LY"
    ]
}

(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ aws eks update-kubeconfig --name eks-g85pk5LY --region us-west-1
Added new context arn:aws:eks:us-west-1:838240021663:cluster/eks-g85pk5LY to C:\
Users\YuriiiHordiichuk\.kube\config
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ kubectl create clusterrolebinding spark-rolebinding \
    --clusterrole=edit \
    --serviceaccount=default:spark
clusterrolebinding.rbac.authorization.k8s.io/spark-rolebinding created
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ kubectl create serviceaccount spark
serviceaccount/spark created
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main
$ eksctl create iamserviceaccount --name spark --namespace default --cluster eks
-g85pk5LY --attach-policy-arn arn:aws:iam::838240021663:policy/S3SparkAccessPol
icy --approve --override-existing-serviceaccounts
2025-11-17 15:48:55 [!] 1 iamserviceaccount (default/spark) was included (based
on the include/exclude rules)
2025-11-17 15:48:55 [!] metadata of serviceaccounts that exist in Kubernetes wi
ll be updated, as --override-existing-serviceaccounts was set
2025-11-17 15:48:55 [!] 1 task: {
    2 sequential sub-tasks: {
        create IAM role for serviceaccount "default/spark",
        create serviceaccount "default/spark",
    }
}
2025-11-17 15:48:55 [!] building iamserviceaccount stack "eksctl-eks-g85
pk5LY-addon-iamserviceaccount-default-spark"
2025-11-17 15:48:55 [!] deploying stack "eksctl-eks-g85pk5LY-addon-iamserviceac
count-default-spark"
2025-11-17 15:48:55 [!] waiting for CloudFormation stack "eksctl-eks-g85pk5LY-a
ddon-iamserviceaccount-default-spark"
2025-11-17 15:49:26 [!] waiting for CloudFormation stack "eksctl-eks-g85pk5LY-a
ddon-iamserviceaccount-default-spark"
2025-11-17 15:49:28 [!] serviceaccount "default/spark" already exists
2025-11-17 15:49:28 [!] updated serviceaccount "default/spark"
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main/spark_config
$ kubectl apply -f spark-role.yaml
role.rbac.authorization.k8s.io/spark-role created
(.venv)
AzureAD+YuriiiHordiichuk@EPPLWROW0218 MINGW64 ~/Desktop/m06_sparkbasics_python_aw
s-main/m06_sparkbasics_python_aws-main/spark_config
$ kubectl apply -f spark-rolebinding.yaml
rolebinding.rbac.authorization.k8s.io/spark-rolebinding created
(.venv)
```

spark-submit run:

```
Administrator: Windows PowerShell
PS C:\WINDOWS\system32> spark-submit `>>   --master "k8s://https://6FBAT74FD2D0B83737D4CCB24983F0832.gr7.us-west-1.eks.amazonaws.com" `>>   --deploy-mode cluster `>>   --name sparkbasics `>>   --conf spark.kubernetes.container.image=838240021663.dkr.ecr.us-west-1.amazonaws.com/spark-python-06:latest `>>   --conf spark.kubernetes.driver.request.cores=500m `>>   --conf spark.kubernetes.driver.request.memory=2g `>>   --conf spark.kubernetes.executor.request.memory=1g `>>   --conf spark.kubernetes.executor.request.cores=500m `>>   --conf spark.kubernetes.namespace=default `>>   --conf spark.kubernetes.authenticate.driver.serviceAccountName=spark `>>   --conf spark.kubernetes.executor.instances=1 `>>   --conf spark.pyspark.python=/usr/bin/python3 `>>   --conf spark.pyspark.driver.python=/usr/bin/python3 `>>   --py-files local:///opt/sparkbasics-1.0.0-py3.11.egg `>>   local:///opt/src/main/python/main_etl.py`> 25/11/17 16:09:44 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable`> 25/11/17 16:09:45 INFO SparkKubernetesClientFactory: Auto-configuring K8S client using current context from users K8S config file`> 25/11/17 16:09:46 INFO KerberosConfDriverFeatureStep: You have not specified a krb5.conf file locally or via a ConfigMap. Make sure that you have the krb5.conf locally on the driver image.`> 25/11/17 16:09:49 INFO LoggingPodStatusWatcherImpl: State changed, new state:`>     pod name: sparkbasics-58230d9a925d62bd-driver`>     namespace: default`>     labels: spark-app-name -> sparkbasics, spark-app-selector -> spark-5c2eacb5aa154a7393a384e8dd9b1f23, spark-role -> driver, spark-version -> 3.5.7`>     pod uid: d96019bd-79dc-43cd-b6ae-1634e502dd3b`>     creation time: 2025-11-17T15:09:47Z`>     service account name: spark`>     volumes: aws-iam-token, spark-local-dir-1, spark-conf-volume-driver, kube-api-access-zh2jc`>     node name: ip-10-0-3-137.us-west-1.compute.internal`>     start time: 2025-11-17T15:09:47Z`>     phase: Pending`>     container status:`>       container name: spark-kubernetes-driver`>       container image: 838240021663.dkr.ecr.us-west-1.amazonaws.com/spark-python-06:latest
```

```
Administrator: Windows PowerShell
    container started at: 2025-11-17T15:09:49Z
    container finished at: 2025-11-17T15:18:23Z
    exit code: 0
    termination reason: Completed
25/11/17 16:18:24 INFO LoggingPodStatusWatcherImpl: Application status for spark-5c2eacb5aa154a7393a384e8dd9b1f23 (phase: Running)
25/11/17 16:18:25 INFO LoggingPodStatusWatcherImpl: State changed, new state:
    pod name: sparkbasics-58230d9a925d62bd-driver
    namespace: default
    labels: spark-app-name -> sparkbasics, spark-app-selector -> spark-5c2eacb5aa154a7393a384e8dd9b1f23, spark-role -> driver, spark-version -> 3.5.7
    pod uid: d96019bd-79dc-43cd-b6ae-1634e502dd3b
    creation time: 2025-11-17T15:09:47Z
    service account name: spark
    volumes: aws-iam-token, spark-local-dir-1, spark-conf-volume-driver, kube-api-access-zh2jc
    node name: ip-10-0-3-137.us-west-1.compute.internal
    start time: 2025-11-17T15:09:47Z
    phase: Succeeded
    container status:
        container name: spark-kubernetes-driver
        container image: 838240021663.dkr.ecr.us-west-1.amazonaws.com/spark-python-06:latest
        container state: terminated
        container started at: 2025-11-17T15:09:49Z
        container finished at: 2025-11-17T15:18:23Z
        exit code: 0
        termination reason: Completed
25/11/17 16:18:25 INFO LoggingPodStatusWatcherImpl: Application status for spark-5c2eacb5aa154a7393a384e8dd9b1f23 (phase: Succeeded)
25/11/17 16:18:25 INFO LoggingPodStatusWatcherImpl: Container final statuses:

    container name: spark-kubernetes-driver
    container image: 838240021663.dkr.ecr.us-west-1.amazonaws.com/spark-python-06:latest
    container state: terminated
    container started at: 2025-11-17T15:09:49Z
    container finished at: 2025-11-17T15:18:23Z
    exit code: 0
    termination reason: Completed
25/11/17 16:18:25 INFO LoggingPodStatusWatcherImpl: Application sparkbasics with application ID spark-5c2eacb5aa154a7393a384e8dd9b1f23 and submission ID default:sparkbasics-58230d9a925d62bd-driver finished
25/11/17 16:18:25 INFO ShutdownHookManager: Shutdown hook called
25/11/17 16:18:25 INFO ShutdownHookManager: Deleting directory C:\Users\YuriHordiichuk\AppData\Local\Temp\spark-2cf92429-0efb-4f33-8ba4-02aa787b9be7
```

Cluster:

eks-g85pk5LY

End of standard support for Kubernetes version 1.32 is March 23, 2026.

Cluster info

Status: Active	Kubernetes version: 1.32	Support period: Standard support until March 23, 2026	Provider: EKS
Cluster health: 0	Upgrade insights: 4 pending, 1 critical	Node health issues: 0	

Overview | Resources | **Compute** | Networking | Add-ons | Access | Observability | Update history & backups | Tags

Nodes (6)

Node name	Instance type	Compute	Managed by	Created	Status
ip-10-0-1-43.us-west-1.compute.internal	t3.medium	Node group	node-group-1-20251117141254223700000019	an hour ago	Ready
ip-10-0-2-108.us-west-1.compute.internal	t3.medium	Node group	node-group-2-2025111714125422430000001b	an hour ago	Ready
ip-10-0-2-139.us-west-1.compute.internal	t3.medium	Node group	node-group-1-20251117141254223700000019	an hour ago	Ready
ip-10-0-2-64.us-west-1.compute.internal	t3.medium	Node group	node-group-1-20251117141254223700000019	an hour ago	Ready
ip-10-0-3-137.us-west-1.compute.internal	t3.medium	Node group	node-group-2-2025111714125422430000001b	an hour ago	Ready

⚠️ Amazon EKS will no longer publish EKS-optimized Amazon Linux 2 (AL2) AMIs after November 26th, 2025. Additionally, Kubernetes version 1.32 is the last version for which Amazon EKS will release AL2 AMIs. From version 1.33 onwards, Amazon EKS will continue to release AL2023 and Bottlerocket based AMIs. [Learn more](#)

Node groups (2)

Group name	Desired size	AMI release version	Launch template	Status
node-group-1-20251117141254223700000019	3	1.32.9-20251112	one-20251117141246027400000017 (1)	Active
node-group-2-2025111714125422430000001b	3	1.32.9-20251112	two-20251117141246024200000015 (1)	Active

Results:

Amazon S3 > Buckets > m06awsbucket-g85pk5Ly > m06sparkbasics/ > result/

result/

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permission.

Name	Type	Last modified	Size	Storage class
SUCCESS	-	November 17, 2025, 17:17:22 (UTC+01:00)	0 B	Standard
part-00000-1bb67209-39b6-4e42-88ae-9ce96eb63cc8-c000.snapy.parquet	parquet	November 17, 2025, 17:17:21 (UTC+01:00)	729.1 KB	Standard

result_df:

```
✓ result_s3_path = "s3a://m06awsbucket-g85pk5ly/m06sparkbasics/result/"  
df_result = spark.read.option("recursiveFileLookup", "true").parquet(result_s3_path)  
✓ 3.4s  
  
df_result.show(10)  
✓ 2.3s  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
|geohash|Id|Name|Country|City|Address|Latitude|Longitude|avg_tmpr_f|avg_tmpr_c|  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
| 9yqg|1477468749829|PhICmvZwu97R5N27y...|zBs9dL9DhsW5CmJSn...|yebLLCCWOYRGK13v...|WlRpjjeBrjeXPv/SK...| 35.79068|-91.63828|71.40808757819484| 21.89186773905272|  
| 9zwx| 592705486848|UYMPj+p+uI2000rVK...|rGxtAxtIprvK0d6JL...|CE4kKh3La11xv/k1t...|/x11Fj+2KWAwPB2+u...| 43.56077|-92.010704|60.20548007246376|15.669429347826076|  
| dp88| 1142461300736|J60TlgjoX4S925/XP...|cI/xkN+KC0+uS8mLs...|gaqYgtnt2D0uoWRnwg...|ThIInyFrQjCgjG0o...| 42.27388|-88.9672|63.04345598509549|17.247275267815553|  
| dr3b| 1073741824001|bjBmf76za0+HjycK...|rQIkMaOVpmbvIk6...|riFeDLik/7xOpVxm8...|oyJmiGF5/8voifouy...| 40.814148|-76.13433|60.52094685127289| 15.84452880750335|  
| gcpv| 2052994367489|7ncapd58s5YKRbLy...|OQV9yCmeuBgbdf3Qx...|q4B+KYIEapNbbs40h...|yv/pkjefnPpkCwTok...| 51.5066991|-0.0280208|56.02065217391304|13.345652173913045|  
| drv5| 506886140929|nk90CNRvYbP8RJzEY...|0+tZ5r0hd/MIXaUa...|gTczDht16tOLVTVq...|x0/uApR2bdr9k+b1j...| 44.266968|-71.56098|54.13619660882683| 12.29834634818558|  
| gcpv| 2671469658114|IsQoB3rZ991AuyfQU...|hU1tuIk9Beu+PDaqH...|AnheLA+rIu6arTgeg...|T12UqIp61rTXMR30L...| 51.5065578|-0.004514|56.02065217391304|13.345652173913045|  
| gcpv| 2748779069440|ylpixKFza4XygLeai...|37LuEN54A31PmRahm...|ZT9ey6554+MF5u4uM...|JYmZqNj+00PFObry/...| 51.5426054|-0.0072708|56.02065217391304|13.345652173913045|  
| sp3e| 2972117368834|Ix487rT/mbhv+7lJ...|gKRp0sFz3Kt0R9eUB...|X199V5eGe4uuJxQbH...|fSmjn4GxG/v+OYjCm...| 41.3944029| 2.150194|70.22282608695652|21.233152173913044|  
| dpc0| 773094113281|Sibfw19x90l3obP3u...|kAhd1BsCw2U7gB1...|/aGsBAuWCE80SuNT8...|ze6EdBno8Q0M007ps...| 43.75521|-88.47649|60.44180253623188| 15.80108695652174|  
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
only showing top 10 rows  
  
df_result.count()  
✓ 1.5s  
2494
```